

Contact

+9145298996 (Mobile)
nileshbhagat1488@gmail.com

www.linkedin.com/in/nilesh-bhagat-data (LinkedIn)

Top Skills

Retail Banking
Health Insurance
Healthcare

Certifications

Vector Databases for RAG: An Introduction
Generative AI: Prompt Engineering Basics
Build RAG Applications: Get Started
Develop Generative AI Applications: Get Started
Generative AI Fundamentals

Nilesh Bhagat

Senior Specialist Data Engineer at LTM | Actuarial & Healthcare Analytics | Azure, Databricks, Snowflake, PySpark, ADF, Big Data | GenAI & RAG | 12+ yrs building cloud data platforms | Open to relocate

Nagpur, Maharashtra, India

Summary

Actively open to Senior / Lead Data Engineer roles in the UAE, Saudi Arabia, and Qatar — visa sponsorship welcome. 12+ years building Azure + Databricks + Snowflake data platforms for global enterprises (banking, insurance, telecom). Let's connect: nileshbhagat1488@gmail.com / +91 9145298996

Senior Specialist Data Engineer with 12+ years of experience designing and delivering cloud-native data platforms on Azure, Databricks, and Snowflake for global enterprises. I build the data foundations that power analytics, ML, and Generative AI at scale. What I do best:

- Design scalable ETL/ELT pipelines with Azure Data Factory, Databricks (PySpark), and Snowflake — handling GBs to multi-TB workloads
- Architect Delta Lake / Lakehouse / Medallion solutions with Unity Catalog governance
- Build AI-ready data pipelines: chunking, embeddings, vector stores (Chroma, FAISS, Milvus) for RAG use cases
- Implement restartable, observable pipelines with Azure Monitor, Logic Apps, and granular audit logging
- Optimize performance — recently cut a claims pipeline from 4 hours to 1 hour using memory-optimized clusters + Delta Cache
- CI/CD for data with Azure DevOps

Recent measurable wins:

- Reduced pipeline runtime by 75% through compute right-sizing and Liquid Clustering
- Built restartable multi-source pipelines (Snowflake + SQL Server + JSON), saving compute cost and SLA breaches
- Enabled RAG-ready datasets for enterprise GenAI applications

Core stack: Azure • Databricks • Snowflake • PySpark • Python • SQL • ADF • Delta Lake • Unity Catalog • Azure DevOps • RAG • Vector Databases • LangChain • LlamaIndex • Lakehouse • Medallion Architecture

Experience

LTIMindtree
Senior Specialist-Data Engineer

July 2024 - Present (2 years)

Pune District

Project: Modern Actuarial Data Warehouse

- 1) Architect and build scalable, cloud-native data pipelines using Databricks, PySpark, ADF, Apache Spark, and Kafka to ingest medical, member, pharmacy, dental, and risk score data from Artemis for actuarial teams.
- 2) Designed and implemented end-to-end Databricks pipelines using Delta Live Tables, Autoloader, and Structured Streaming, reducing data ingestion latency by 45% across multiple production workloads.
- 3) Developed structured and unstructured data ingestion from multiple sources including APIs, databases, and streaming data sources.
- 4) Built automated pipelines to generate actuarial reports by aggregating claims, premium, and membership data from Artemis with full audit trail.
- 5) Designed and implemented RAG pipelines enabling enterprise knowledge retrieval using embedding generation, vector databases (ChromaDB), and LLM integration.
- 6) Built and enforced data quality frameworks using DQX, implementing validation rules, SLA monitoring, and automated alerting across 15+ pipelines — achieving 99.2% data reliability in production.
- 7) Collaborated directly with actuarial and underwriting teams to translate reporting requirements into scalable pipeline designs on Delta Lake.

UST

Data Engineer

September 2021 - August 2024 (3 years)

Pune, Maharashtra, India

Project: Loan Scoring System

- 1) Engineered a scalable loan scoring system by integrating diverse financial data sources into a unified data lake architecture using Azure Data Factory and PySpark.

- 2) Built automated ETL pipelines to clean, transform, and standardize credit, payment, and transactional data, enabling accurate and timely loan risk assessment.
- 3) Collaborated closely with data scientists to provide feature-ready datasets for ML scoring models, improving loan approval decision speed and reducing defaults.
- 4) Applied robust data validation and profiling logic to ensure high-quality input for the scoring engine, directly improving model reliability.
- 5) Supported Power BI integration for real-time loan risk dashboards used by underwriters and finance executives.
- 6) Enhanced operational efficiency by automating recurring data loads and reducing manual intervention by 70%.

Birlasoft

Senior Software Engineer

January 2020 - September 2021 (1 year 9 months)

Pune, Maharashtra, India

Developing and maintaining data pipelines for the extraction, transformation, and loading (ETL) of data from various sources to the data warehouse.

Ensuring data quality and integrity throughout the ETL process

Facilitated POCs for Data migration to Azure, delivering 100% efficiency.

Capita

Senior Software Engineer

July 2018 - January 2020 (1 year 7 months)

Pune Area, India

Assisted in designing and implementing ETL processes to move and transform data from source systems to data warehouses.

Contributed to the design and implementation of data models, schemas, and structures

Learn and contribute to the automation of data processes and workflows to improve efficiency.

Cognizant

Programming Analyst

August 2016 - July 2018 (2 years)

Pune, Maharashtra, India

Tested ETL processes for concurrency and parallel processing to ensure they can handle multiple data loads simultaneously

Improved test coverage through the implementation of comprehensive test plans and the integration of Gherkin-based BDD (Behavior-Driven Development) testing.

Implemented test automation using Selenium, resulting in a reduction in annual testing costs.

Infosys

Technology Specialist

July 2014 - January 2016 (1 year 7 months)

Pune Area, India

Validates data by running queries in the database and verifying that the results match the expectations

Automation Test cases execution and logs checking

Validate the accuracy of metadata, including table structures, data types, and relationships.

ETH Limited

Software Developer

March 2013 - July 2014 (1 year 5 months)

Pune, Maharashtra, India

Developed web-based applications written in Java

Identify and fix bugs in the code.

Work closely with team members, including other developers, designers, and product managers.

Education

Nagpur University

Bachelor of Engineering - BE, Information Technology · (2007 - 2012)